

Comparing regression and LSTM sentiment analysis models

Language Processing II exam
Jens Egholm Pedersen
xtp778@sc.ku.dk

June 16, 2017

Contents

1	Introduction	3
2	Background	3
2.1	Language processing	3
2.1.1	Sentiment analysis	4
2.2	Linear regression	4
2.2.1	Error reporting	4
2.3	Neural networks	5
2.3.1	Long short-term memory networks	5
2.3.2	Optimisation and regularization	6
2.3.3	Error reporting	7
3	Data generation	7
3.1	Learning models	7
3.1.1	Data set	7
3.1.2	K-fold cross validation	7
4	Quantitative evaluation	7
5	Qualitative evaluation	7
5.1	Other approaches	8
6	Conclusion	8
7	References	9
8	Appendix A: Data	10
9	Appendix B: Software	11
9.1	NLTK	11
9.2	Scikit-learn	11
9.3	Tensorflow	11
9.4	Keras	11

10 Appendix C: Figure sources	12
10.1 Figure 4	12

1 Introduction

Due to the enormous rise in accessible data and processing capabilities, computers have been moving from the restricted formalized domain of mathematics and computability, and into the more complex and seemingly chaotic domain of natural language (Nilsson, 2009; Jurafsky and Martin, 2000). Automating understanding of natural language have a wide range of helpful applications, ranging from translation to recommender systems to generic personal assistants (Cox, 2005; Bird et al., 2009).

As a part of this development, sentiment analysis has been attracting increased attention (Bird et al., 2009; Jurafsky and Martin, 2000).

In the present paper, sentiment analysis is applied to a corpus of movie reviews presented by Pang and Lee (2005). By using standard natural language processing tools and machine learning techniques, four models are built and trained to predict the rating of a movie review on a discrete normalized scale.

2 Background

This section illuminates the theory and background for language processing and basic methodology for processing text. Sentiment analysis as a fundamental concept for this setting will be introduced second, followed by fundamental statistical metrics. Lastly, recurrent neural networks will be introduced before moving on to the data generation.

2.1 Language processing

Following the revolutionizing formalizations of grammar by Chomsky (2002), first written in 1957, much work was put into working with the computability of language (Jurafsky and Martin, 2000). In the beginning of the 21st century, the previous work converged with the statistical machine learning community, giving rise to supervised, semi-supervised and even unsupervised models of language understanding and translation (Jurafsky and Martin, 2000).

There are several approaches to language processing, but it is common to deconstruct text into its constituents and assign meaning to each piece, with the hope that the disassembly gives additional meaning that can aggregate to an understanding of the text as a whole when the pieces are put back together (Jurafsky and Martin, 2000). *Constituents* can range from morphemes to words to whole sentences in this context. This process is referred to as feature extraction due to the enrichment of the text with meta-data.

A fundamental method for feature extraction is to construct a probabilistic model by simply joining items (ranging from phonemes to word pairs) together in pairs of one, two or more, called n -grams (Jurafsky and Martin, 2000). The advantage of this approach focuses on speed and the simple, yet efficient, probabilities in for instance word pairs which can be applied to predict likelihood of future occurrences (Jurafsky and Martin, 2000).

Term frequency, inverse document frequency (tf-idf) is another relevant feature that describes the relevance of an item, weighted against the number of times it appears in other documents (Jurafsky and Martin, 2000) (see figure 1). Tf-idf favours words that is used heavily in a single document (assumed to be

$$tf - idf = Nw \cdot \log\left(\frac{N}{k}\right)$$

Figure 1: Formula for tf-idf.

important) while diminishing the importance of words that often occur in other documents, and is a heavily used metric (Jurafsky and Martin, 2000).

2.1.1 Sentiment analysis

Another popular approach for disassembling text called part of speech-tagging (POS tagging), assigns grammatical categories to each word and attempt to build a linguistic model that can be understood by a machine (Jurafsky and Martin, 2000). This works well for information retrieval, where objective truths can be extracted¹, but fails to capture the complexity and emotional variety of sentences that appear outside the grammar, such as emotions and opinions² (Jurafsky and Martin, 2000).

Sentiment analysis focuses on extracting information about the emotions or opinions of a text (Pang and Lee, 2008). Defining what a *sentiment* is not a simple task, and have been scrutinized extensively (Jurafsky and Martin, 2000; Pang and Lee, 2008). A much used approach was introduced by Ekman (1992) where he divided six basic emotions, which could be conveniently operationalized by computers. Since the dataset for this paper revolves around opinions in a single dimension, *sentiment* will in this paper simply refer to a subjective experience, discretized to a number (see section 3.1.1).

2.2 Linear regression

Linear regression is a mathematical function that describes how the mean of the output changes with the value of the input (Agresti and Finlay, 2008). The model is appropriate when a linear correlation between the input and output can be assumed, so an increase in the input results in a linear in- or decrease in the output, represented by a straight line in a 2-dimensional coordinate system (Agresti and Finlay, 2008).

2.2.1 Error reporting

The predictions made by linear regression model (y), does not always fit the actual values (\hat{y}). To describe this error and to give indicators on the quality of the model, the mean absolute error (MAE) and root mean squared error (RMSE) is introduced (Agresti and Finlay, 2008; Chai and Draxler, 2014).

MAE in figure 2 describes the mean of the absolute error values. If a model has an MAE of 1, the average distance to the actual value from the predicted value is 1. The bigger the mean error value, the worse the model is.

¹Such as the objective statements about an object in the sentence "*The cat likes tuna fish*".

²One example is a heavily sarcastic sentence like "*I'd really truly love going out in this weather!*". For POS-tagging sarcasm is difficult to capture because it is not implicitly encoded in the grammatical categories.

$$MAE = \frac{\sum |e|}{n}, e = y - \hat{y}$$

Figure 2: Mean absolute error (MAE).

$$RMSE = \sqrt{\frac{\sum (e)^2}{n-2}}, e = y - \hat{y}$$

Figure 3: Root mean square error (RMSE).

However, this figure does not capture the error distribution of the model (Chai and Draxler, 2014). If a gaussian distribution is assumed for the errors, it is important to know the standard distribution to indicate *how* consistently wrong the model is (Agresti and Finlay, 2008). The root mean square error measures the standard deviation in the sample of the errors, making it a good description for the model accuracy, compared to other models (Agresti and Finlay, 2008; Chai and Draxler, 2014).

Each metric should not stand alone however, and MAE and RMSE is typically reported together (Chai and Draxler, 2014).

2.3 Neural networks

Inspired by the biological brain, models for neural networks have been applied in computer science the 1950s, and have become integral part of machine learning (Nilsson, 2009; Russell and Norvig, 2009). Neurons are essentially functions that operate on some input and respond with some output (Russell and Norvig, 2009). The inputs for a neuron are weighted to give different input channels - or input dimensions - varying significance. These weighted inputs arrive to an activation function that decides whether the neuron should ‘fire’ or not (Nilsson, 2009). Sigmoid or hyperbolic tangent are popular activation functions for numerical predictions because of their steep logistic properties while retaining differentiability. Neural networks excell by their adaptability and have long been applied to the field of language processing (Jurafsky and Martin, 2000).

Simple neural networks can be stacked in layers, where each neuron will assign weights to the input, to determine its significance for the activation function (Nilsson, 2009). By tuning the weights of the neurons, such groupings can learn to fire on certain input patterns (Russell and Norvig, 2009). However, this structure have proved brittle and difficult to train because the neurons do not retain their weights for long (Nilsson, 2009; Russell and Norvig, 2009). Rumelhart et al. (1988) suggested a method to avoid this instability by adjusting the weights of the neurons in retrospect with a method called back-propagation (Rumelhart et al., 1988; Nilsson, 2009). This optimization is operationalized by a function that can describe the *loss* of efficiency in a network, and then adjust the weights correspondingly using gradient descent (Russell and Norvig, 2009).

2.3.1 Long short-term memory networks

Back-propagating networks with a high complexity in layers and neurons have shown to be hard to optimize above a certain point, because the gradient de-

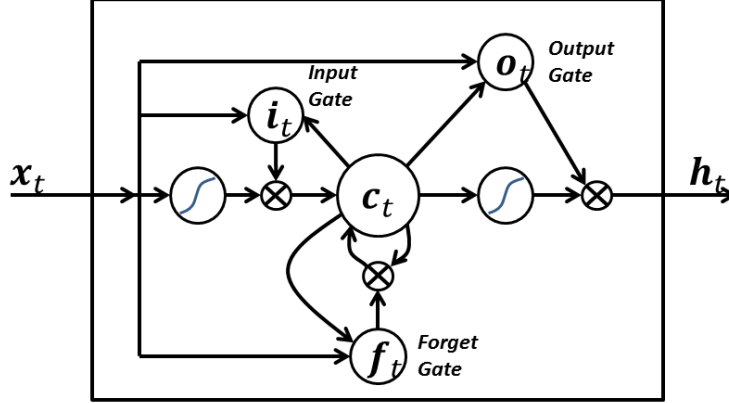


Figure 4: An *peephole* LSTM unit with three gates (input, output and forget).

scent algorithm are inefficient at escaping local extrema in a high-dimensional space (Russell and Norvig, 2009). Recurrence have been used to counter this problem, by conserving or returning to previous, optimal values (Schmidhuber, 2015; Russell and Norvig, 2009). Hochreiter and Schmidhuber (1997) suggested a certain type of memory-retaining networks that is able to balance between both storing previous optimal state, while also forgetting context when new and different input arrives (Hochreiter and Schmidhuber, 1997; Schmidhuber, 2015). They dubbed the model long short-term memory (LSTM) networks, to capture both properties (Hochreiter and Schmidhuber, 1997).

The LSTM model seen in figure 4 builds on the idea of a more complex neuronal unit, with several *gates*, collaborating to give the LSTM the desired property (Gers, 2001). A traditional unit contains three gates: an "input" gate that control the degree with which new information should influence the current memory state, a "forget" gate that decides whether a value should be retained and an output gate that decides how much influence the memory should have on the activation of the unit (Hochreiter and Schmidhuber, 1997; Gers, 2001). There are a plentitude of other architectures, but for this paper, the LSTM model shown in figure 4 is used (Gers, 2001).

2.3.2 Optimisation and regularization

Recurrent neural networks are typically trained with a continuously applied loss function and gradiend descent for weight correction (?Hochreiter and Schmidhuber, 1997; ?). Stochastic gradiens descent and resilient back-propagation are two common approaches, which have proved efficient (?Nilsson, 2009). This paper applies a third method, based on the RMSE algorithm in figure 3, called root mean square propagation (RMSprop), which preserves the stochastic properties while controlling for previous errors in the network (Tieleman and Hinton, 2012).

Model	r^2	MAE	RMSE
Linear	0	0.1649	0.2075
Linear + tf-idf	0	0.1494	0.1855
LSTM	N/A	0.1667	0.2045
LSTM + tf-idf	N/A	0.1641	0.2014

Figure 5: Comparison of metrics for the four different models using mean absolute error (MAE) and root mean squared error (RMSE).

2.3.3 Error reporting

3 Data generation

This section

3.1 Learning models

The four models model is built using the Scikit Learn (Sklearn) package and second is based on the deep learning library Keras (see appendix in section 9). Two The corpus is evaluated against four different implementations of sentiment analysis, based on two different models. The first model uses the

3.1.1 Data set

Chose to focus on ratings

3.1.2 K-fold cross validation

The dataset

N-grams Jurafsky and Martin (2000).

max features for bow (10'000)

this section must report on actual test-runs utilizing the movie review corpus and applying at least two machine learning algorithms (Naïve-Bayes being one example)

4 Quantitative evaluation

demonstrating how SA results can be evaluated automatically

5 Qualitative evaluation

evaluating and discussing the quantitative results wrt. validity, reliability and/or relevance in a “real -world” perspective

The very definition of *sentiment* in section ?? as a subjective opinion makes comparison difficult. As stated in (Pang and Lee, 2005) the entries in the present dataset is difficult to compare upon, since ratings are relative between reviewers; a rating of 0.8 might be high for one reviewer, but low for another. Unfortunately a simple normalization does not solve the problem because the scale of one reviewer might not even be linearly comparable to another.

5.1 Other approaches

Not taken AFINN, leksika word2vec Nielsen (2011)

6 Conclusion

presenting in a condensed form your results and observations, e.g. pointing to strengths, weaknesses, and future directions

7 References

- Agresti, A. and Finlay, B. (2008). *Statistical Methods for the Social Sciences (4th Edition)*. Prentice Hall, 4 edition.
- Bird, S., Klein, E., and Loper, E. (2009). *Natural Language Processing with Python: Analyzing Text with the Natural Language Toolkit*. O'Reilly, Beijing.
- Chai, T. and Draxler, R. R. (2014). Root mean square error (RMSE) or mean absolute error (MAE)? - Arguments against avoiding RMSE in the literature. *Geoscientific Model Development*, 7:1247–1250.
- Chomsky, N. (2002). *Syntactic Structures*. Mouton classic. Bod Third Party Titles.
- Cox, M. T. (2005). Metacognition in computation: A selected research review. *Artificial Intelligence*, 169(2):104 – 141. Special Review Issue.
- Ekman, P. (1992). An argument for basic emotions. *Cognition and Emotion*, pages 169–200.
- Gers, F. (2001). Long short-term memory in recurrent neural networks.
- Hochreiter, S. and Schmidhuber, J. (1997). Long short-term memory. *Neural Comput.*, 9(8):1735–1780.
- Jurafsky, D. and Martin, J. H. (2000). *Speech and Language Processing: An Introduction to Natural Language Processing, Computational Linguistics, and Speech Recognition*. Prentice Hall PTR, Upper Saddle River, NJ, USA, 1st edition.
- Nielsen, F. Å. (2011). AFINN.
- Nilsson, N. J. (2009). *The quest for artificial intelligence - A history of ideas and achievements*. Cambridge University Press.
- Pang, B. and Lee, L. (2005). Seeing stars: Exploiting class relationships for sentiment categorization with respect to rating scales. In *Proceedings of ACL*, pages 115–124.
- Pang, B. and Lee, L. (2008). Opinion mining and sentiment analysis. *Foundations and Trends in Information Retrieval*, 2:1–135.
- Rumelhart, D. E., Hinton, G. E., and Williams, R. J. (1988). Neurocomputing: Foundations of research. chapter Learning Representations by Back-propagating Errors, pages 696–699. MIT Press, Cambridge, MA, USA.
- Russell, S. and Norvig, P. (2009). *Artificial Intelligence: A Modern Approach*. Prentice Hall Press, Upper Saddle River, NJ, USA, 3rd edition.
- Schmidhuber, J. (2015). Deep learning in neural networks: An overview. *Neural Networks*, 61:85 – 117.
- Tieleman, T. and Hinton, G. (2012). Lecture 6.5-rmsprop: Divide the gradient by a running average of its recent magnitude.

8 Appendix A: Data

The data used in this report is collected from Pang and Lee (2005). The data is available in the software repository (see 9). The readme with further descriptions on the data source is available at:

<http://www.cs.cornell.edu/people/pabo/movie-review-data/scaledata.README.1.0.txt>

9 Appendix B: Software

All software used to produce the results presented in this report is open-source and available through GitHub at:

<https://github.com/Jegp/langprocexam>

Below follows a list of the software used. An in-depth description on how to reproduce the results are available from the GitHub repository above.

9.1 NLTK

A toolkit for natural language processing in python. Visited 15th of June 2017.

<http://www.nltk.org/api/nltk.sentiment.html>

9.2 Scikit-learn

Machine learning library for Python. Visited 15th of June 2017.

<http://scikit-learn.org/stable/index.html>

9.3 Tensorflow

A machine learning library initially developed by Google Inc. Visited 15th of June 2017.

<https://www.tensorflow.org/>

9.4 Keras

A deep learning library which builds on top of Tensorflow. Visited 4th of June 2017.

<https://keras.io>

10 Appendix C: Figure sources

10.1 Figure 4

Source: Wikipedia.org

https://en.wikipedia.org/wiki/File:Peephole_Long_Short-Term_Memory.svg.